

Convolutional Neural Networks for Image Processing

Image Processing with Neural Networks:

images as data: `data.shape` \rightarrow (^{height} 2832, ^{number of channels (RGB)} 4256, 3)

`data[250, 3500]` \rightarrow array (^{width} [0.74, 0.078, 0.145]) \rightarrow it's a red pixel!
_{row idx column idx R G B}

input and modify images:

```
import matplotlib.pyplot as plt
```

```
data = plt.imread('image.jpg')
```

```
data[:, :, 1] = 0  $\rightarrow$  set green channel as zero
```

```
data[:, :, 2] = 0  $\rightarrow$  set blue channel as zero
```

```
plt.imshow(data)  $\rightarrow$  input image with only RED channel.  
plt.show()
```

change an image: `data[200:1200, 200:1200, :] = [0, 1, 0]` \rightarrow make it green!

in black & white images \rightarrow brighter pixels have higher values \rightarrow 255 white
 \rightarrow 0 black

represent categorical image labels mathematically: One-hot Encoding

```
categories = np.array(["t-shirt", "dress", "shoe"])
```

```
n_categories = 3
```

```
one_labels = np.zeros((len(labels), n_categories))
```

```
for ii in range(len(labels)):
```

```
    jj = np.where(categories == labels[ii])
```

```
    one_labels[ii, jj] = 1
```

★ One-hot encoding can be used for test predictions

test array ($\begin{bmatrix} 1, 0, 0 \\ 0, 1, 0 \\ \vdots \end{bmatrix}$, prediction ($\begin{bmatrix} 1, 0, 0 \\ 0, 0, 1 \\ \vdots \end{bmatrix}$, \Rightarrow (test * prediction).sum()
 \downarrow
num of correct classifications

Keras for image classification:

From keras.models import Sequential

```
model = Sequential()
```

from keras.layers import Dense

model.add(Dense(10, activation = 'relu', input_shape = (784,)))

```
model.add(Dense(10, activation s 'relu'))
```

```
model.add(Dense(3, activation='softmax'))
```

model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

```
train_data = train_data.reshape((50, 784))
```

model.fit(train_data, train_labels, validation_split=0.2, epochs=3)
to prevent overfitting → validation

→ training: model adjusts weights through backpropagation and gradient descent

```
test_data = test_data.reshape((10, 784))
```


```
model.evaluate(test_data, test_labels)
```

number of images $C(BW)$

pe $\rightarrow (50, 28, 28, 1)$

$\underbrace{\hspace{1.5cm}}_h \quad \underbrace{\hspace{1.5cm}}_w$

28 x 28



→ go over all data 3 times

multi-class

to prevent overfitting \rightarrow validation

Using Convolutions:

Natural images have spatial correlations \rightarrow e.g. pixels along a contour or edge


```
array = np.array([0,0,0,0,0,1,1,1])
```

kernel = np.array([-1, 1]) \rightarrow 1-D Convolution \rightarrow image: 2D Convolution

```
Conv = np.array([0, ..., 0])
```

kernel = $\begin{bmatrix} -1 & 1 \end{bmatrix} \Rightarrow$ Left Edge Detector!

for i in range(8):

$$\text{Conv}(ii) = (\text{kernel} * \text{array}[ii; ii+2]).\text{sum}() \rightarrow \text{Conv} \Rightarrow \text{array}([0, 0, 0, 0, 1, 0])$$


Kernel: $\begin{bmatrix} -1 & 1 & -1 \\ -1 & 1 & -1 \\ -1 & 1 & -1 \end{bmatrix} \rightarrow$ Finds vertical lines

$$\begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix} \rightarrow \text{horizontal lines}$$
$$\begin{bmatrix} -1 & -1 & -1 \\ -1 & 1 & -1 \\ -1 & -1 & -1 \end{bmatrix} \rightarrow \text{light spot in the dark}$$
$$\begin{bmatrix} 1 & 1 & 1 \\ 1 & -1 & 1 \\ 1 & 1 & 1 \end{bmatrix} \rightarrow \text{dark spot in the light}$$

From keras.layers import Conv2D

padding is the same as Dense layer but with fewer weights

Conv2D(10, kernel_size=3, activation='relu')

↳ number of filters ↳ 3x3

↳ Convolution has one weight for each pixel in the kernel

From keras.models import Sequential

num parameters = $10 \times 3 \times 3 = 90$

From keras.layers import Dense, Conv2D, Flatten

model = Sequential()

model.add(Conv2D(10, kernel_size=3, activation='relu', input_shape=(img_rows, img_columns, 1)))

num channels

model.add(Flatten()) → connection between Convolution and Dense layer

↳ Flatten the Conv output into a one dimensional array

model.add(Dense(3, activation='Softmax'))

↳ three classes

model.compile(...)

→ here we do not reshape input data, because we want to have the spatial relationships → so we should define input shape in the Conv2D

model.fit(..., batch_size=10)

↳ also

→ cross entropy loss

model.evaluate(...) → output: [0.547, 1.0] → accuracy!

zero padding: add zeros to base image to have same size for the feature map after convolution

model.add(Conv2D(..., padding='valid') → default

padding='same' → zero padding applied!

Stride: size of the step we take with the kernel over the input image

model.add(Conv2D(..., strides=1) → strides > 1 → output image is smaller than input

input size

kernel size

size of zero padding

calculate the size of output:

$$O = \frac{I - K + 2P}{S} + 1$$

↳ S → strides

★ Dilated Convolution → useful when we want to aggregate information across multiple scales

model.add(Conv2D(..., dilation=2)

Going Deeper:

input_shape keyword argument is required for first layer only. ~~Other~~ other layers do not need this argument to be specified.

~~Deep~~ Deep Network → more than one convolutional layer

Why adding more layers?


Features in early layers → simple textures
Features in intermediate layers → more complex features (simple objects)
Features in late layers → specific types of objects

Deep Networks → more computational cost
→ may require more data

hierarchical representation of data

How many parameters in a model? → model.summary()


description of model + total trainable and non-trainable parameters



Fully Connected model

$$784 \times 10 + 10 \quad 10 \times 10 + 10 \quad 3 \times 10 + 3 \rightarrow 7993 \text{ (trainable) parameters}$$

bias terms



Convolution model
(padding = 'same', stride = 1)

$$9 \times 10 + 10 \quad 9 \times 10 \times 10 + 10 \quad 28 \times 28 \times 10 \times 3 + 3 \rightarrow 24533 \text{ (trainable) parameters}$$

Kernel values (3x3) no. kernels bias term

★ Convolution networks do not necessarily reduce number of parameters. In FC networks most number of parameters are for the first dense layer. In contrast, in CNNs most parameters come from the ~~last layers~~ last layers. (mainly the first dense layer after the flattening)

pooling operations: Convolve on image with a kernel. on each point, replace the related part of the image with the maximum value available on that part. (max pooling)
 if the kernel is $2 \times 2 \rightarrow$ the image dimensions reduce to a quarter after pooling.

```
result = np.zeros((im.shape[0] // 2, im.shape[1] // 2))
result[0,0] = np.max(im[0:2, 0:2])
result[0,1] = np.max(im[0:2, 2:4])
...
```

\Downarrow in Keras

From Keras.layers import MaxPool2D

model = Sequential()

model.add(Conv2D(...))

model.add(MaxPool2D(2))

model.add(Conv2D(...))

model.add(MaxPool2D(2))

...

put pooling layer after each Conv layer


kernel size

model.summary()

\downarrow
 number of parameters reduced from ~ 30000 to ~ 2000 !

Understanding and Improving Deep CNNs

is learning progressive as expected? \rightarrow view changes in the loss during learning



training = model.fit(train_data, train_labels, ...)

import matplotlib.pyplot as plt

\rightarrow "training" has a dictionary to store the learning curves

plt.plot(training.history['loss'])

plt.plot(training.history['val_loss'])

plt.show()


contains functions which can be executed at the end of each training epoch.

★ Storing the optimal parameters before the model starts overfitting?

```
from keras.callbacks import ModelCheckpoint  → stores the weights
checkpoint = ModelCheckpoint('weights.hdf5', monitor='val_loss',
                             file_type='h5', save_best_only=True)
callbacks_list = [checkpoint]  it only gets lists
model.fit(train_data, train_labels, ..., callbacks=callbacks_list)
model.load_weights('weights.hdf5')
model.predict_classes(test_data) → array([2, 1, 2, 0, ...])
```

Neural network regularization: to prevent overfitting

→ Dropout: In each learning step, select a subset of units and ignore them in the forward pass, and, in the back-propagation of error



→ introduced in 2014

```
from keras.layers import Dropout
```

```
⋮
```

```
model.add(Conv2D(...))
```

```
model.add(Dropout(0.25))  → proportion of the units to be neglected
                           (randomly changes the selected units in each learning step)
```

```
⋮
```

→ Batch Normalization: Rescale the outputs of each layer → in 2015

```
from keras.layers import BatchNormalization
```

```
⋮
```

```
model.add(Conv2D(...))
```

```
model.add(BatchNormalization())
```

```
⋮
```

→ Be careful! Sometimes models become worse when using these two together.

The disharmony between batch normalization and dropout

Interpreting the model: CNNs are black boxes! → it will evolve rapidly, soon!

model.layers → [<... Conv>, <... Conv>, <... Flatten>, <dense>]


conv1 = model.layers[0]

weights1 = conv1.get_weights() → len(weights) = 2

kernels1 = weights1[0] → kernels1.shape = (3, 3, 1, 9)

num_channels
num_kernels
3x3 kernels

kernels1_1 = kernels1[:, :, 0, 0] → kernels1_1.shape = (3, 3)

plt.imshow(kernels1_1) →  → but, what it does?!

test_image = test_data[3, :, :, 0] ← convolve that with one image

filtered_image = convolution(test_image, kernels1_1) → we coded the function before! (few pages back)

plt.imshow(filtered_image)

Further reading : <https://distill.pub/2017/Feature-visualization/>
by Chris Olah →

